

Nouman Ejaz

Bahria Town, Lahore, Pakistan | +92 324 1428934 | working.nouman.ejaz@gmail.com

PROFESSIONAL SUMMARY

Architecting and scaling enterprise-grade AI platforms powering millions of users across real-time and high-throughput environments. As an AI/ML and Cloud Engineer with a strong full-stack foundation, I design and ship autonomous AI agent systems, RAG pipelines, and ML-powered products. Proven ability to architect cloud-native infrastructure (AWS, GCP, Azure), optimize large-scale LLM inference systems, and lead cross-functional engineering teams. Dedicated to building scalable, well-structured systems that adhere to SOLID principles and deliver measurable business value.

TECHNICAL SKILLS

Languages: Python, JavaScript, TypeScript, C++, R, Rust, Golang, SQL

AI & Agentic Systems: LangChain, LangGraph, CrewAI, AutoGen, Agent2Agent Protocol, LlamaIndex, Semantic Kernel, MCP, OpenAI Agents SDK

ML & Data Science: PyTorch, TensorFlow, scikit-learn, Hugging Face Transformers, NumPy, Pandas, Jupyter

MLOps & API Integration: MLflow, Weights & Biases, OpenAI API, Claude API, Gemini API, DeepSeek, Qwen

Vector & Graph Databases: Pinecone, Qdrant, Azure Cosmos DB, ChromaDB, Neo4j, Amazon Neptune, Memgraph

Cloud & Infrastructure: AWS, Google Cloud (GCP), Azure, Terraform, Docker, Kubernetes, GitHub Actions

Frontend & Backend Dev: React JS, Next JS, Node.js, Express JS, FastAPI, GraphQL, Socket.IO

Databases & ORMs: PostgreSQL, MySQL, MongoDB, Redis, Prisma, Sequelize, Mongoose, SQLite

PROFESSIONAL EXPERIENCE

Lead AI/ML Engineer

Rizviz Internation Impex

Feb 2026 – Present

Lahore, Pakistan

- Architected and delivered an enterprise-scale RAG platform integrating LLMs, vector databases, external APIs, and structured data to drastically improve knowledge retrieval.
- Designed and deployed multi-agent AI systems using LangGraph and AutoGen, enabling autonomous workflows and end-to-end process automation across various enterprise use cases.
- Engineered a high-throughput LLM inference platform supporting millions of daily requests; implemented intelligent batching, caching, and cost-aware routing to significantly reduce infrastructure costs.
- Optimized inference performance via token-level optimization and distributed serving, sharply improving latency under heavy production loads.
- Built a robust LLM evaluation framework combining offline benchmarking and online A/B testing to guarantee model reliability, consistency, and production trustworthiness.
- Led and mentored a cross-functional team of 12 engineers across ML, platform, and infrastructure teams, driving architectural strategy, roadmap planning, and alignment.

Associate Software Engineer (AI Specialist)

GenIteam Solutions

Aug 2025 – Mar 2026

Lahore, Pakistan

- Designed and scaled distributed ML pipelines supporting both batch and real-time inference workloads at production scale.
- Collaborated directly with international clients and stakeholders, including NVIDIA, to develop and deliver a complex research simulation project.
- Built and deployed recommender systems and fraud detection models, driving improved model performance that directly impacted core business KPIs.

- Architected a centralized feature store and reusable training pipelines, accelerating the team's experimentation lifecycle and model iteration speed.
- Developed automated monitoring and retraining pipelines to ensure continuous model performance and proactively handle data drift in production.

Full Stack Associate Software Engineer

Feb 2025 – Jul 2025

Onyxtec Solution

Lahore, Pakistan

- Built large-scale, full-stack systems for enterprise clients, specifically developing real-time applications directly for the healthcare industry.
- Designed and deployed highly scalable, REST-based microservices to ensure real-time self-healing and rigorous system compliance.
- Enhanced machine learning model accuracy through advanced feature engineering and comprehensive hyperparameter optimization.

Senior Frontend Engineer

Jul 2024 – Feb 2025

CareerTweakz Inc.

Lahore, Pakistan

- Developed highly interactive, large-scale frontend applications for AI startups and enterprise clients.
- Designed complex analytical dashboards to support business intelligence, data visualization, and executive decision-making.
- Automated internal data workflows, significantly improving system reliability and operational efficiency across organizational data platforms.

WordPress Developer & Automation Engineer

Jan 2023 – Mar 2024

ExceloWeb Solutions

Remote

- Developed customized e-commerce platforms and online stores for US-based clients in a fully remote capacity.
- Authored Python scripts to automate routine business processes, saving substantial manual hours.
- Engineered web scraping tools to extract and process research and marketing data, driving strategic business improvements.

AI SYSTEM AND ARCHITECTURE HIGHLIGHTS

- **Enterprise RAG:** Designed a large-scale Retrieval-Augmented Generation system integrating diverse documents, APIs, and structured data, reducing knowledge retrieval time by 70% while drastically improving response quality.
- **Multi-Agent AI:** Built an advanced agentic system capable of autonomous planning, complex reasoning, and tool execution, featuring human-in-the-loop validation to ensure safe and reliable enterprise automation.
- **Real-Time Fraud:** Architected a high-throughput streaming pipeline utilizing Kafka and Spark, achieving sub-second inference latency for mission-critical security workloads.
- **LLM Optimization:** Implemented sophisticated batching, semantic caching, and intelligent routing strategies to maintain high-scale performance while aggressively minimizing compute costs.

EDUCATION

Bachelor of Science in Software Engineering

2021 – 2025

University of Management and Technology

Lahore, Pakistan